

Running Citus using Docker

The smallest possible Citus cluster is a single PostgreSQL node with the Citus extension, which means you can try out Citus by running a single Docker container.

```
# run PostgreSQL with Citus on port 5500
docker run -d --name citus -p 5500:5432 -e POSTGRES_PASSWORD=mypassword
citusdata/citus

# connect using psql within the Docker container
docker exec -it citus psql -U postgres

# or, connect using local psql
psql -U postgres -d postgres -h localhost -p 5500
```

Install Citus locally

If you already have a local PostgreSQL installation, the easiest way to install Citus is to use our packaging repo

Install packages on Ubuntu / Debian:

```
curl https://install.citusdata.com/community/deb.sh > add-citus-repo.sh
sudo bash add-citus-repo.sh
sudo apt-get -y install postgresql-16-citus-12.1
```

Install packages on CentOS / Red Hat:

```
curl https://install.citusdata.com/community/rpm.sh > add-citus-repo.sh
sudo bash add-citus-repo.sh
sudo yum install -y citus121_16
```

To add Citus to your local PostgreSQL database, add the following to `postgresql.conf`:

```
shared_preload_libraries = 'citus'
```

After restarting PostgreSQL, connect using `psql` and run:

```
CREATE EXTENSION citus;
```

You're now ready to get started and use Citus tables on a single node.

Install Citus on multiple nodes

If you want to set up a multi-node cluster, you can also set up additional PostgreSQL nodes with the Citus extensions and add them to form a Citus cluster:

```
-- before adding the first worker node, tell future worker nodes how to reach
the coordinator
```

```
SELECT citus_set_coordinator_host('10.0.0.1', 5432);

-- add worker nodes
SELECT citus_add_node('10.0.0.2', 5432);
SELECT citus_add_node('10.0.0.3', 5432);

-- rebalance the shards over the new worker nodes
SELECT rebalance_table_shards();
```

For more details, see our [documentation on how to set up a multi-node Citus cluster](#) on various operating systems.

Using Citus

Once you have your Citus cluster, you can start creating distributed tables, reference tables and use columnar storage.

Creating Distributed Tables

The `create_distributed_table` UDF will transparently shard your table locally or across the worker nodes:

```
CREATE TABLE events (
    device_id bigint,
    event_id bigserial,
    event_time timestamptz default now(),
    data jsonb not null,
    PRIMARY KEY (device_id, event_id)
);

-- distribute the events table across shards placed locally or on the worker nodes
SELECT create_distributed_table('events', 'device_id');
```

After this operation, queries for a specific device ID will be efficiently routed to a single worker node, while queries across device IDs will be parallelized across the cluster.

```
-- insert some events
INSERT INTO events (device_id, data)
SELECT s % 100, ('{"measurement":'||random()||'}')::jsonb FROM
generate_series(1,1000000) s;

-- get the last 3 events for device 1, routed to a single node
SELECT * FROM events WHERE device_id = 1 ORDER BY event_time DESC, event_id
DESC LIMIT 3;
```

device_id	event_id	event_time	data

1	1999901	2021-03-04 16:00:31.189963+00	{"measurement": 0.88722643925054}
1	1999801	2021-03-04 16:00:31.189963+00	{"measurement": 0.6512231304621992}
1	1999701	2021-03-04 16:00:31.189963+00	{"measurement": 0.019368766051897524}

(3 rows)

Time: 4.588 ms

-- explain plan for a query that is parallelized across shards, which shows the plan for

-- a query one of the shards and how the aggregation across shards is done
EXPLAIN (VERBOSE ON) SELECT count(*) FROM events;

```

QUERY PLAN

Aggregate
  Output: COALESCE((pg_catalog.sum(remote_scan.count))::bigint, '0'::bigint)
  -> Custom Scan (Citrus Adaptive)
    ...
    -> Task
      Query: SELECT count(*) AS count FROM events_102008 events
      Node: host=localhost port=5432 dbname=postgres
      -> Aggregate
        -> Seq Scan on public.events_102008 events
WHERE true

```

Creating Distributed Tables with Co-location

Distributed tables that have the same distribution column can be co-located to enable high performance distributed joins and foreign keys between distributed tables. By default, distributed tables will be co-located based on the type of the distribution column, but you define co-location explicitly with the `colocate_with` argument in `create_distributed_table`.

```

CREATE TABLE devices (
  device_id bigint primary key,
  device_name text,
  device_type_id int
);

```

```

CREATE INDEX ON devices (device_type_id);

-- co-locate the devices table with the events table
SELECT create_distributed_table('devices', 'device_id', colocate_with :=
'events');

-- insert device metadata
INSERT INTO devices (device_id, device_name, device_type_id)
SELECT s, 'device-'||s, 55 FROM generate_series(0, 99) s;

-- optionally: make sure the application can only insert events for a known
device
ALTER TABLE events ADD CONSTRAINT device_id_fk
FOREIGN KEY (device_id) REFERENCES devices (device_id);

-- get the average measurement across all devices of type 55, parallelized
across shards
SELECT avg((data->>'measurement')::double precision)
FROM events JOIN devices USING (device_id)
WHERE device_type_id = 55;

```

avg
0.5000191877513974

(1 row)

Time: 209.961 ms

Co-location also helps you scale [INSERT..SELECT](#), [stored procedures](#), and [distributed transactions](#).

Distributing Tables without interrupting the application

Some of you already start with Postgres, and decide to distribute tables later on while your application using the tables. In that case, you want to avoid downtime for both reads and writes. `create_distributed_table` command block writes (e.g., DML commands) on the table until the command is finished. Instead, with `create_distributed_table_concurrently` command, your application can continue to read and write the data even during the command.

```

CREATE TABLE device_logs (
    device_id bigint primary key,
    log text
);

-- insert device logs
INSERT INTO device_logs (device_id, log)
SELECT s, 'device log:'||s FROM generate_series(0, 99) s;

-- convert device_logs into a distributed table without interrupting the
application
SELECT create_distributed_table_concurrently('device_logs', 'device_id',
colocate_with := 'devices');

```

```
-- get the count of the logs, parallelized across shards
SELECT count(*) FROM device_logs;
```

count
100

(1 row)

Time: 48.734 ms

Creating Reference Tables

When you need fast joins or foreign keys that do not include the distribution column, you can use `create_reference_table` to replicate a table across all nodes in the cluster.

```
CREATE TABLE device_types (
  device_type_id int primary key,
  device_type_name text not null unique
);
```

```
-- replicate the table across all nodes to enable foreign keys and joins on
any column
```

```
SELECT create_reference_table('device_types');
```

```
-- insert a device type
```

```
INSERT INTO device_types (device_type_id, device_type_name) VALUES (55,
'laptop');
```

```
-- optionally: make sure the application can only insert devices with known
types
```

```
ALTER TABLE devices ADD CONSTRAINT device_type_fk
FOREIGN KEY (device_type_id) REFERENCES device_types (device_type_id);
```

```
-- get the last 3 events for devices whose type name starts with laptop,
parallelized across shards
```

```
SELECT device_id, event_time, data->>'measurement' AS value, device_name,
device_type_name
FROM events JOIN devices USING (device_id) JOIN device_types USING
(device_type_id)
WHERE device_type_name LIKE 'laptop%' ORDER BY event_time DESC LIMIT 3;
```

device_id	event_time	value	device_name	device_type_name
60	2021-03-04 16:00:31.189963+00	0.28902084163415864	laptop	
8	2021-03-04 16:00:31.189963+00	0.8723803076285073	laptop	

```

      20 | 2021-03-04 16:00:31.189963+00 | 0.8177634801548557 | device-20
laptop |
(3 rows)

```

Time: 146.063 ms

Reference tables enable you to scale out complex data models and take full advantage of relational database features.

Creating Tables with Columnar Storage

To use columnar storage in your PostgreSQL database, all you need to do is add `USING columnar` to your `CREATE TABLE` statements and your data will be automatically compressed using the columnar access method.

```

CREATE TABLE events_columnar (
    device_id bigint,
    event_id bigserial,
    event_time timestamptz default now(),
    data jsonb not null
)
USING columnar;

-- insert some data
INSERT INTO events_columnar (device_id, data)
SELECT d, '{"hello":"columnar"}' FROM generate_series(1,10000000) d;

-- create a row-based table to compare
CREATE TABLE events_row AS SELECT * FROM events_columnar;

-- see the huge size difference!
\d+

```

List of relations					
Schema	Name	Type	Owner	Persistence	
public	events_columnar	table	marco	permanent	25 MB
public	events_row	table	marco	permanent	651 MB

(2 rows)

You can use columnar storage by itself, or in a distributed table to combine the benefits of compression and the distributed query engine.

When using columnar storage, you should only load data in batch using `COPY` or `INSERT..SELECT` to achieve good compression. Update, delete, and foreign keys are currently unsupported on columnar tables. However, you can use partitioned tables in which newer partitions use row-based storage, and older partitions are compressed using columnar storage.

To learn more about columnar storage, check out the [columnar storage README](#).

Schema-based sharding

Available since Citus 12.0, [schema-based sharding](#) is the shared database, separate schema model, the schema becomes the logical shard within the database. Multi-tenant apps can use a schema per tenant to easily shard along the tenant dimension. Query changes are not required and the application usually only needs a small modification to set the proper `search_path` when switching tenants. Schema-based sharding is an ideal solution for microservices, and for ISVs deploying applications that cannot undergo the changes required to onboard row-based sharding.

Creating distributed schemas

You can turn an existing schema into a distributed schema by calling `citus_schema_distribute`:

```
SELECT citus_schema_distribute('user_service');
```

Alternatively, you can set `citus.enable_schema_based_sharding` to have all newly created schemas be automatically converted into distributed schemas:

```
SET citus.enable_schema_based_sharding TO ON;
```

```
CREATE SCHEMA AUTHORIZATION user_service;  
CREATE SCHEMA AUTHORIZATION time_service;  
CREATE SCHEMA AUTHORIZATION ping_service;
```

Running queries

Queries will be properly routed to schemas based on `search_path` or by explicitly using the schema name in the query.

For [microservices](#) you would create a `USER` per service matching the schema name, hence the default `search_path` would contain the schema name. When connected the user queries would be automatically routed and no changes to the microservice would be required.

```
CREATE USER user_service;  
CREATE SCHEMA AUTHORIZATION user_service;
```

For typical multi-tenant applications, you would set the search path to the tenant schema name in your application:

```
SET search_path = tenant_name, public;
```

Setting up with High Availability

One of the most popular high availability solutions for PostgreSQL, [Patroni 3.0](#), has [first class support for Citus 10.0 and above](#), additionally since Citus 11.2 ships with improvements for smoother node switchover in Patroni.

An example of patronictl list output for the Citus cluster:

```
postgres@coord1:~$ patronictl list demo
+ Citus cluster: demo -----+-----+-----+-----+
| Group | Member | Host           | Role           | State   | TL | Lag in MB |
+-----+-----+-----+-----+-----+-----+-----+
| 0      | coord1  | 172.27.0.10    | Replica        | running | 1  | 0          |
| 0      | coord2  | 172.27.0.6     | Sync Standby   | running | 1  | 0          |
| 0      | coord3  | 172.27.0.4     | Leader         | running | 1  |            |
| 1      | work1-1 | 172.27.0.8     | Sync Standby   | running | 1  | 0          |
| 1      | work1-2 | 172.27.0.2     | Leader         | running | 1  |            |
| 2      | work2-1 | 172.27.0.5     | Sync Standby   | running | 1  | 0          |
| 2      | work2-2 | 172.27.0.7     | Leader         | running | 1  |            |
+-----+-----+-----+-----+-----+-----+-----+
```

Documentation

If you're ready to get started with Citus or want to know more, we recommend reading the [Citus open source documentation](#). Or, if you are using Citus on Azure, then the [Azure Cosmos DB for PostgreSQL](#) is the place to start.

Our Citus docs contain comprehensive use case guides on how to build a [multi-tenant SaaS application](#), [real-time analytics dashboard](#), or work with [time series data](#).

Architecture

A Citus database cluster grows from a single PostgreSQL node into a cluster by adding worker nodes. In a Citus cluster, the original node to which the application connects is referred to as the coordinator node. The Citus coordinator contains both the metadata of distributed tables and reference tables, as well as regular (local) tables, sequences, and other database objects (e.g. foreign tables).

Data in distributed tables is stored in “shards”, which are actually just regular PostgreSQL tables on the worker nodes. When querying a distributed table on the coordinator node, Citus will send regular SQL queries to the worker nodes. That way, all the usual PostgreSQL optimizations and extensions can automatically be used with Citus.

When you send a query in which all (co-located) distributed tables have the same filter on the distribution column, Citus will automatically detect that and send the whole query to the worker node that stores the data. That way, arbitrarily complex queries are supported with minimal routing overhead, which is especially useful for scaling transactional workloads. If queries do not have a specific filter, each shard is queried in parallel, which is especially useful in analytical workloads. The Citus distributed executor is adaptive and is designed to handle both query types at the same time on the same system under high concurrency, which enables large-scale mixed workloads.

The schema and metadata of distributed tables and reference tables are automatically synchronized to all the nodes in the cluster. That way, you can connect to any node to run distributed queries. Schema changes and cluster administration still need to go through the coordinator.

Detailed descriptions of the implementation for Citus developers are provided in the [Citus Technical Documentation](#).

When to use Citus

Citus is uniquely capable of scaling both analytical and transactional workloads with up to petabytes of data. Use cases in which Citus is commonly used:

- **[Customer-facing analytics dashboards](#)**: Citus enables you to build analytics dashboards that simultaneously ingest and process large amounts of data in the database and give sub-second response times even with a large number of concurrent users.

The advanced parallel, distributed query engine in Citus combined with PostgreSQL features such as [array types](#), [JSONB](#), [lateral joins](#), and extensions like [HyperLogLog](#) and [TopN](#) allow you to build responsive analytics dashboards no matter how many customers or how much data you have.

Example real-time analytics users: [Algolia](#)

- **[Time series data](#)**: Citus enables you to process and analyze very large amounts of time series data. The biggest Citus clusters store well over a petabyte of time series data and ingest terabytes per day.

Citus integrates seamlessly with [Postgres table partitioning](#) and has [built-in functions for partitioning by time](#), which can speed up queries and writes on time series tables. You can take advantage of Citus's parallel, distributed query engine for fast analytical queries, and use the built-in *columnar storage* to compress old partitions.

Example users: [MixRank](#)

- **[Software-as-a-service \(SaaS\) applications](#)**: SaaS and other multi-tenant applications need to be able to scale their database as the number of tenants/customers grows. Citus

enables you to transparently shard a complex data model by the tenant dimension, so your database can grow along with your business.

By distributing tables along a tenant ID column and co-locating data for the same tenant, Citus can horizontally scale complex (tenant-scoped) queries, transactions, and foreign key graphs. Reference tables and distributed DDL commands make database management a breeze compared to manual sharding. On top of that, you have a built-in distributed query engine for doing cross-tenant analytics inside the database.

Example multi-tenant SaaS users: [Salesloft](#), [ConvertFlow](#)

- **[Microservices](#)**: Citus supports schema based sharding, which allows distributing regular database schemas across many machines. This sharding methodology fits nicely with typical Microservices architecture, where storage is fully owned by the service hence can't share the same schema definition with other tenants. Citus allows distributing horizontally scalable state across services, solving one of the [main problems](#) of microservices.
- **Geospatial**: Because of the powerful [PostGIS](#) extension to Postgres that adds support for geographic objects into Postgres, many people run spatial/GIS applications on top of Postgres. And since spatial location information has become part of our daily life, well, there are more geospatial applications than ever. When your Postgres database needs to scale out to handle an increased workload, Citus is a good fit.