Non-Relational Postgres
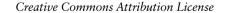
BRUCE MOMJIAN

EDB

This talk explores the advantages of non-relational storage, and the Postgres support for such storage.

Relational Storage

- Relational storage was proposed by E. F. Codd in 1970
- 50+ years still in use
- Very flexible
- Not always ideal

https://www.youtube.com/watch?v=zSn8il5Mo5s

What Is Relational Storage?

- Row, column, table (tuple, attribute, relation)
- Constraints
- **Normalization, joins**

- Each column/attribute contains only atomic indivisible values
- Eliminate repeating groups in individual tables
- Create a separate table for each set of related data
- Identify each set of related data with a primary key

https://en.wikipedia.org/wiki/First_normal_form
http://www.anaesthetist.com/mnm/sql/normal.htm

Downsides of First Normal Form

- Query performance
- Query complexity
- Storage inflexibility
- Storage overhead
- Indexing limitations

Postgres Non-Relational Storage Options

1. Arrays
2. Range types
3. Geometry
4. Xml
5. Json
6. Jsonb
7. Row types
8. Character strings

1. Arrays

```
CREATE TABLE employee
(name TEXT PRIMARY KEY, certifications TEXT[]);

INSERT INTO employee
VALUES ('Bill', '{"CCNA", "ACSP", "CISSP"}');

SELECT * FROM employee;
 name |   certifications
------+-------------------
 Bill | {CCNA,ACSP,CISSP}

SELECT name
FROM employee
WHERE certifications @> '{ACSP}';
 name
------
 Bill
```

All queries used in this presentation are available at https://momjian.us/main/writings/pgsql/
non-relational.sql.

Array Access

```
SELECT certifications[1]
FROM employee;
 certifications
----------------
 CCNA

SELECT unnest(certifications)
FROM employee;
 unnest
--------
 CCNA
 ACSP
 CISSP
```

Array Unrolling

```
SELECT name, unnest(certifications)
FROM employee;
 name | unnest
------+--------
 Bill | CCNA
 Bill | ACSP
 Bill | CISSP
```

Array Creation

```
SELECT DISTINCT relkind
FROM pg_class
ORDER BY 1;
 relkind
---------
 i
 r
 t
 v

SELECT array_agg(DISTINCT relkind)
FROM pg_class;
 array_agg
-----------
 {i,r,t,v}
```

```
CREATE TABLE car_rental
(id SERIAL PRIMARY KEY, time_span TSTZRANGE);

INSERT INTO car_rental
VALUES (DEFAULT, '[2016-05-03 09:00:00, 2016-05-11 12:00:00)');

SELECT *
FROM car_rental
WHERE time_span @> '2016-05-09 00:00:00'::timestamptz;
 id |                    time_span
----+----------------------------------------------------
  1 | ["2016-05-03 09:00:00-04","2016-05-11 12:00:00-04")

SELECT *
FROM car_rental
WHERE time_span @> '2018-06-09 00:00:00'::timestamptz;
 id | time_span
----+-----------
```

Range Type Indexing

```
INSERT INTO car_rental (time_span)
SELECT tstzrange(y, y + '1 day')
FROM generate_series('2001-09-01 00:00:00'::timestamptz,
                     '2010-09-01 00:00:00'::timestamptz, '1 day') AS x(y);

SELECT *
FROM car_rental
WHERE time_span @> '2007-08-01 00:00:00'::timestamptz;
  id |                     time_span
------+-------------------------------------------------------
 2162 | ["2007-08-01 00:00:00-04","2007-08-02 00:00:00-04")

EXPLAIN SELECT *
FROM car_rental
WHERE time_span @> '2007-08-01 00:00:00'::timestamptz;
                                  QUERY PLAN
--------------------------------------------------------------------------------
 Seq Scan on car_rental  (cost=0.00..64.69 rows=16 width=36)
   Filter: (time_span @> '2007-08-01 00:00:00-04'::timestamp with time zone)
```

12/71

Range Type Indexing

```
CREATE INDEX car_rental_idx ON car_rental
USING GIST (time_span);

EXPLAIN SELECT *
FROM car_rental
WHERE time_span @> '2007-08-01 00:00:00'::timestamptz;
                                    QUERY PLAN
-----------------------------------------------------------------------------------------
 Bitmap Heap Scan on car_rental  (cost=4.27..28.35 rows=16 width=36)
   Recheck Cond: (time_span @> '2007-08-01 00:00:00-04'::timestamp with time zone)
   ->  Bitmap Index Scan on car_rental_idx  (cost=0.00..4.27 rows=16 width=0)
         Index Cond: (time_span @> '2007-08-01 00:00:00-04'::timestamp with time zone)
```

Exclusion Constraints

```
ALTER TABLE car_rental ADD EXCLUDE USING GIST (time_span WITH &&);

INSERT INTO car_rental
VALUES (DEFAULT, '[2003-04-01 00:00:00, 2003-04-01 00:00:01)');
ERROR:  conflicting key value violates exclusion constraint "car_rental_time_span_excl"
DETAIL:   Key (time_span)=(["2003-04-01 00:00:00-05","2003-04-01 00:00:01-05"))
conflicts with existing key (time_span)=(["2003-04-01 00:00:00-05","2003-04-02 00:00:00-05")).
```

3. Geometry

```
CREATE TABLE dart (dartno SERIAL, location POINT);


INSERT INTO dart (location)
SELECT CAST('(' || random() * 100 || ',' ||
            random() * 100 || ')' AS point)
FROM generate_series(1, 1000);


SELECT *
FROM dart
LIMIT 5;
 dartno |              location
--------+-------------------------------------
      1 | (60.1593657396734,64.1712633892894)
      2 | (22.9252253193408,38.7973457109183)
      3 | (54.7123382799327,16.1387695930898)
      4 | (60.5669556651264,53.1596980988979)
      5 | (22.7800350170583,90.8143546432257)
```

Geometry Restriction

```
-- find all darts within four units of point (50, 50)
SELECT *
FROM dart
WHERE location <@ '<(50, 50), 4>'::circle;
 dartno |               location
--------+------------------------------------
    308 | (52.3920683190227,49.3803130928427)
    369 | (52.1113255061209,52.9995835851878)
    466 | (47.5943599361926,49.0266934968531)
    589 | (46.3589935097843,50.3238912206143)
    793 | (47.3468563519418,50.0582652166486)

EXPLAIN SELECT *
FROM dart
WHERE location <@ '<(50, 50), 4>'::circle;
                       QUERY PLAN
--------------------------------------------------------
 Seq Scan on dart  (cost=0.00..19.50 rows=1 width=20)
   Filter: (location <@ '<(50,50),4>'::circle)
```

Indexed Geometry Restriction

```
CREATE INDEX dart_idx ON dart
USING GIST (location);

EXPLAIN SELECT *
FROM dart
WHERE location <@ '<(50, 50), 4>'::circle;
                                QUERY PLAN
-------------------------------------------------------------------
 Index Scan using dart_idx on dart  (cost=0.14..8.16 rows=1 width=20)
   Index Cond: (location <@ '<(50,50),4>'::circle)
```

```
-- find the two closest darts to (50, 50)
SELECT *
FROM dart
ORDER BY location <-> '(50, 50)'::point
LIMIT 2;
 dartno |              location
--------+------------------------------------
    308 | (52.3920683190227,49.3803130928427)
    466 | (47.5943599361926,49.0266934968531)

EXPLAIN SELECT *
FROM dart
ORDER BY location <-> '(50, 50)'::point
LIMIT 2;
                              QUERY PLAN
--------------------------------------------------------------------------------
 Limit  (cost=0.14..0.30 rows=2 width=28)
   -> Index Scan using darts_idx on darts  (cost=0.14..80.14 rows=1000 width=28)
         Order By: (location <-> '(50,50)'::point)
```

```
$ # Run with foomatic installed, or download:
$ # https://www.openprinting.org/download/foomatic/foomatic-db-4.0-current.tar.gz
$ cd /usr/share/foomatic/db/source/opt
$ for FILE in *.xml
do    tr -d '\n' < "$FILE"
      echo
done > /tmp/foomatic.xml

$ psql
CREATE TABLE printer (doc XML);

COPY printer from '/tmp/foomatic.xml';
```

Xpath Query

```
SELECT (xpath('/option/arg_shortname/en', doc))
FROM printer
LIMIT 5;
                xpath
--------------------------------
 {<en>Dithering</en>}
 {<en>BottomMargin</en>}
 {<en>Uniform</en>}
 {<en>CurlCorrectionAlways</en>}
 {<en>Encoding</en>}
```

Remove XML Array

```
SELECT (xpath('/option/arg_shortname/en', doc))[1]
FROM printer
LIMIT 5;
              xpath
-------------------------------
 <en>Dithering</en>
 <en>BottomMargin</en>
 <en>Uniform</en>
 <en>CurlCorrectionAlways</en>
 <en>Encoding</en>
```

```
-- convert to XML text
SELECT (xpath('/option/arg_shortname/en/text()', doc))[1]
FROM printer
LIMIT 5;
        xpath
----------------------
 Dithering
 BottomMargin
 Uniform
 CurlCorrectionAlways
 Encoding
```

Xpath to SQL Text

```
-- convert to SQL text so we can do DISTINCT and ORDER BY
SELECT DISTINCT text((xpath('/option/arg_shortname/en/text()', doc))[1])
FROM printer
ORDER BY 1
LIMIT 5;
     text
--------------
 AlignA
 AlignB
 AlignC
 AlignD
 AllowReprint
```

XML Non-Root Query

```
SELECT xpath('//driver/text()', doc)
FROM printer
ORDER BY random()
LIMIT 5;
        xpath
----------------------
 {bjc600,bjc800,hpdj}
 {hl1250}
 {oki182}
 {bjc600,bjc800}
 {Postscript1}
```

Unnest XML Arrays

```
SELECT DISTINCT unnest((xpath('//driver/text()', doc))::text[])
FROM printer
ORDER BY 1
LIMIT 5;
  unnest
----------
 ap3250
 appledmp
 bj10
 bj10e
 bj10v
```

Search XML Text

```
WITH driver (name) AS (
    SELECT DISTINCT unnest(xpath('//driver/text()', doc))::text
    FROM printer
)
SELECT name
FROM driver
WHERE name LIKE 'hp%'
ORDER BY 1;
    name
-------------
 hpdj
 hpijs
 hpijs-pcl3
 hpijs-pcl5c
 hpijs-pcl5e
```

5. JSON Data Type

- JSON data type, not to be confused with JSONB
- Similar to XML in that the JSON is stored as text and validated
- ~100 JSON functions

Load JSON Data

```
-- download sample data from https://www.mockaroo.com/
-- remove 'id' column, output as JSON, uncheck 'array'
CREATE TABLE friend (id SERIAL, data JSON);

COPY friend (data) FROM '/tmp/MOCK_DATA.json';

SELECT *
FROM friend
ORDER BY 1
LIMIT 2;
 id |                            data
----+-----------------------------------------------------------…
  1 | {"gender":"Male","first_name":"Eugene","last_name":"Reed",…
  2 | {"gender":"Female","first_name":"Amanda","last_name":"Morr…
```

Pretty Print JSON

```
SELECT id, jsonb_pretty(data::jsonb)
FROM friend
ORDER BY 1
LIMIT 1;
 id |              jsonb_pretty
----+-----------------------------------------
  1 | {                                       +
    |     "email": "ereed0@businesswire.com",+
    |     "gender": "Male",                   +
    |     "last_name": "Reed",                +
    |     "first_name": "Eugene",             +
    |     "ip_address": "46.168.181.79"      +
    | }
```

```
SELECT data->>'email'
FROM friend
ORDER BY 1
LIMIT 5;
            ?column?
-----------------------------
 aalexandere0@europa.eu
 aalvarezdk@miitbeian.gov.cn
 aandrewsd9@usda.gov
 aarmstrong61@samsung.com
 abarnes55@de.vu
```

Concatenate JSON Values

```
SELECT data->>'first_name' || ' ' ||
       (data->>'last_name')
FROM friend
ORDER BY 1
LIMIT 5;
    ?column?
----------------
 Aaron Alvarez
 Aaron Murphy
 Aaron Rivera
 Aaron Scott
 Adam Armstrong
```

JSON Value Restrictions

```
SELECT data->>'first_name'
FROM friend
WHERE data->>'last_name' = 'Banks'
ORDER BY 1;
 ?column?
----------
 Bruce
 Fred

-- the JSON way
SELECT data->>'first_name'
FROM friend
WHERE data::jsonb @> '{"last_name" : "Banks"}'
ORDER BY 1;
 ?column?
----------
 Bruce
 Fred
```

Single-Key JSON Index

```
-- need double parentheses for the expression index
CREATE INDEX friend_idx ON friend ((data->>'last_name'));

EXPLAIN SELECT data->>'first_name'
FROM friend
WHERE data->>'last_name' = 'Banks'
ORDER BY 1;
                                    QUERY PLAN
-------------------------------------------------------------------------…
 Sort  (cost=12.89..12.90 rows=3 width=123)
   Sort Key: ((data ->> 'first_name'::text))
   -> Bitmap Heap Scan on friend  (cost=4.30..12.87 rows=3 width=123)
        Recheck Cond: ((data ->> 'last_name'::text) = 'Banks'::text)
        -> Bitmap Index Scan on friend_idx  (cost=0.00..4.30 rows=3 …
              Index Cond: ((data ->> 'last_name'::text) = 'Banks'::t…
```

```
SELECT data->>'first_name' || ' ' || (data->>'last_name'),
       data->>'ip_address'
FROM friend
WHERE (data->>'ip_address')::inet <<= '172.0.0.0/8'::cidr
ORDER BY 1;
   ?column?    |     ?column?
---------------+------------------
 Lisa Holmes   | 172.65.223.150
 Walter Miller | 172.254.148.168


SELECT data->>'gender', COUNT(data->>'gender')
FROM friend
GROUP BY 1
ORDER BY 2 DESC;
 ?column? | count
----------+-------
 Male     |   507
 Female   |   493
```

6. JSONB

Like the JSON data type, except:

- Values are native JavaScript data types: text, number, boolean, null, subobject
- Indexing of all keys and values
- Stored in compressed format
- Sorts keys to allow binary-search key look up
- Does not preserve key order
- Does not preserve whitespace syntax
- Retains only the last duplicate key

hstore is similar non-hierarchical key/value implementation.

JSON vs. JSONB Data Types

```
SELECT '{"name" : "Jim", "name" : "Andy", "age" : 12}'::json;
                        json
--------------------------------------------------
 {"name" : "Jim", "name" : "Andy", "age" : 12}

SELECT '{"name" : "Jim", "name" : "Andy", "age" : 12}'::jsonb;
             jsonb
-----------------------------
 {"age": 12, "name": "Andy"}
```

JSONB Index

```
CREATE TABLE friend2 (id SERIAL, data JSONB);

INSERT INTO friend2
SELECT * FROM friend;

-- jsonb_path_ops indexes are smaller and faster,
-- but do not support key-existence lookups.
CREATE INDEX friend2_idx ON friend2
USING GIN (data);
```

JSONB Index Queries

```
SELECT data->>'first_name'
FROM friend2
WHERE data @> '{"last_name" : "Banks"}'
ORDER BY 1;
 ?column?
----------
 Bruce
 Fred

EXPLAIN SELECT data->>'first_name'
FROM friend2
WHERE data @> '{"last_name" : "Banks"}'
ORDER BY 1;                    QUERY PLAN
----------------------------------------------------------------…
 Sort  (cost=24.03..24.04 rows=1 width=139)
   Sort Key: ((data ->> 'first_name'::text))
   -> Bitmap Heap Scan on friend2  (cost=20.01..24.02 rows=1 …
         Recheck Cond: (data @> '{"last_name": "Banks"}'::jsonb)
         -> Bitmap Index Scan on friend2_idx  (cost=0.00..20.01 ……
               Index Cond: (data @> '{"last_name": "Banks"}'::js…
```

JSONB Index Queries

```
SELECT data->>'last_name'
FROM friend2
WHERE data @> '{"first_name" : "Jane"}'
ORDER BY 1;
 ?column?
----------
 Tucker
 Williams

EXPLAIN SELECT data->>'last_name'
FROM friend2
WHERE data::jsonb @> '{"first_name" : "Jane"}'
ORDER BY 1;                    QUERY PLAN
--------------------------------------------------------------…
 Sort  (cost=24.03..24.04 rows=1 width=139)
   Sort Key: ((data ->> 'last_name'::text))
   -> Bitmap Heap Scan on friend2  (cost=20.01..24.02 rows=1 …
         Recheck Cond: (data @> '{"first_name": "Jane"}'::jsonb)
         -> Bitmap Index Scan on friend2_idx  (cost=0.00..20.01 …
               Index Cond: (data @> '{"first_name": "Jane"}'::js…
```

JSONB Index Queries

```
SELECT data->>'first_name' || ' ' || (data->>'last_name')
FROM friend2
WHERE data @> '{"ip_address" : "62.212.235.80"}'
ORDER BY 1;
    ?column?
-----------------
 Theresa Schmidt

EXPLAIN SELECT data->>'first_name' || ' ' || (data->>'last_name')
FROM friend2
WHERE data @> '{"ip_address" : "62.212.235.80"}'
ORDER BY 1;                   QUERY PLAN
-------------------------------------------------------------------…
 Sort  (cost=24.04..24.05 rows=1 width=139)
   Sort Key: ((((data ->> 'first_name'::text) || ' '::text) || …
   -> Bitmap Heap Scan on friend2  (cost=20.01..24.03 rows=1 …
         Recheck Cond: (data @> '{"ip_address": "62.212.235.80"}'…
         -> Bitmap Index Scan on friend2_idx  (cost=0.00..20.01 …
               Index Cond: (data @> '{"ip_address": "62.212.235.…
```

```
CREATE TYPE drivers_license AS
(state CHAR(2), id INTEGER, valid_until DATE);

CREATE TABLE truck_driver
(id SERIAL, name TEXT, license DRIVERS_LICENSE);

INSERT INTO truck_driver
VALUES (DEFAULT, 'Jimbo Biggins', ('PA', 175319, '2017-03-12'));
```

```
SELECT *
FROM truck_driver;
 id |     name     |        license
----+--------------+------------------------
  1 | Jimbo Biggins | (PA,175319,2017-03-12)

SELECT license
FROM truck_driver;
        license
------------------------
 (PA,175319,2017-03-12)

-- parentheses are necessary
SELECT (license).state
FROM truck_driver;
 state
-------
 PA
```

```
$ cd /tmp
$ wget http://web.mit.edu/freebsd/head/games/fortune/datfiles/fortunes
$ psql postgres

CREATE TABLE fortune (line TEXT);

COPY fortune FROM '/tmp/fortunes' WITH (DELIMITER E'\x1F');
```

```
SELECT * FROM fortune WHERE line = 'underdog';
 line
------

SELECT * FROM fortune WHERE line = 'Underdog';
   line
----------
 Underdog

SELECT * FROM fortune WHERE lower(line) = 'underdog';
   line
----------
 Underdog
```

Case Folding

```
CREATE INDEX fortune_idx_text ON fortune (line);

EXPLAIN SELECT * FROM fortune WHERE lower(line) = 'underdog';
                         QUERY PLAN
-------------------------------------------------------------
 Seq Scan on fortune  (cost=0.00..1384.63 rows=295 width=36)
   Filter: (lower(line) = 'underdog'::text)
```

Indexed Case Folding

```
CREATE INDEX fortune_idx_lower ON fortune (lower(line));

EXPLAIN SELECT * FROM fortune WHERE lower(line) = 'underdog';
                                    QUERY PLAN
------------------------------------------------------------…
 Bitmap Heap Scan on fortune  (cost=14.70..468.77 rows=295 …
   Recheck Cond: (lower(line) = 'underdog'::text)
   -> Bitmap Index Scan on fortune_idx_lower  (cost=0.00..…
         Index Cond: (lower(line) = 'underdog'::text)
```

String Prefix

```
SELECT line
FROM fortune
WHERE line LIKE 'Mop%'
ORDER BY 1;
          line
------------------------
 Mophobia, n.:
 Moping, melancholy mad:
```

String Prefix

```
EXPLAIN SELECT line
FROM fortune
WHERE line LIKE 'Mop%'
ORDER BY 1;
                            QUERY PLAN
-----------------------------------------------------------------
 Sort  (cost=1237.07..1237.08 rows=4 width=36)
   Sort Key: line
   -> Seq Scan on fortune  (cost=0.00..1237.03 rows=4 width=36)
         Filter: (line ~~ 'Mop%'::text)
```

Indexed String Prefix

```
-- The default op class does string ordering of non-ASCII
-- collations, but not partial matching.  text_pattern_ops
-- handles prefix matching, but not ordering.
CREATE INDEX fortune_idx_ops ON fortune (line text_pattern_ops);

EXPLAIN SELECT line
FROM fortune
WHERE line LIKE 'Mop%'
ORDER BY 1;
                                    QUERY PLAN
-------------------------------------------------------------------…
 Sort  (cost=8.48..8.49 rows=4 width=36)
   Sort Key: line
   -> Index Only Scan using fortune_idx_ops on fortune  (cost=0.41 …
         Index Cond: ((line ˜>=˜ 'Mop'::text) AND (line ˜<˜ 'Moq'::…
         Filter: (line ˜˜ 'Mop%'::text)
```

Case Folded String Prefix

```
EXPLAIN SELECT line
FROM fortune
WHERE lower(line) LIKE 'mop%'
ORDER BY 1;
                            QUERY PLAN
-----------------------------------------------------------------
 Sort  (cost=1396.73..1397.47 rows=295 width=36)
   Sort Key: line
   -> Seq Scan on fortune  (cost=0.00..1384.63 rows=295 width=36)
        Filter: (lower(line) ~~ 'mop%'::text)
```

Indexed Case Folded String Prefix

```
CREATE INDEX fortune_idx_ops_lower ON fortune
(lower(line) text_pattern_ops);

EXPLAIN SELECT line
FROM fortune
WHERE lower(line) LIKE 'mop%'
ORDER BY 1;
                                    QUERY PLAN
-----------------------------------------------------------------…
 Sort  (cost=481.61..482.35 rows=295 width=36)
   Sort Key: line
   -> Bitmap Heap Scan on fortune  (cost=15.44..469.51 rows=295 …
        Filter: (lower(line) ˜˜ 'mop%'::text)
        -> Bitmap Index Scan on fortune_idx_ops_lower  (cost=0…
              Index Cond: ((lower(line) ˜>=˜ 'mop'::text) AND (…
```

8.2. Full Text Search

- Allows whole-word or word prefix searches
- Supports *and, or, not*
- Converts words to lexemes
    - stemming
    - 21 languages supported
    - `Simple` search config bypasses stemming
- Removes stop words
- Supports synonyms and phrase transformations (thesaurus)

```
SHOW default_text_search_config;
 default_text_search_config
-----------------------------
 pg_catalog.english

SELECT to_tsvector('I can hardly wait.');
    to_tsvector
-------------------
 'hard':3 'wait':4

SELECT to_tsquery('hardly & wait');
   to_tsquery
-----------------
 'hard' & 'wait'
```

```
SELECT to_tsvector('I can hardly wait.') @@
       to_tsquery('hardly & wait');
 ?column?
----------
 t

SELECT to_tsvector('I can hardly wait.') @@
       to_tsquery('softly & wait');
 ?column?
----------
 f
```

```
CREATE INDEX fortune_idx_ts ON fortune
USING GIN (to_tsvector('english', line));
```

Full Text Search Queries

```
SELECT line
FROM fortune
WHERE to_tsvector('english', line) @@ to_tsquery('pandas');
                                  line
-----------------------------------------------------------------------
        A giant panda bear is really a member of the raccoon family.

EXPLAIN SELECT line
FROM fortune
WHERE to_tsvector('english', line) @@ to_tsquery('pandas');
                                    QUERY PLAN
---------------------------------------------------------------------…
 Bitmap Heap Scan on fortune  (cost=12.41..94.25 rows=21 width=36)
   Recheck Cond: (to_tsvector('english'::regconfig, line) @@ to_ts…
   -> Bitmap Index Scan on fortune_idx_ts  (cost=0.00..12.40 rows…
        Index Cond: (to_tsvector('english'::regconfig, line) @@ t…
```

Complex Full Text Search Queries

```
SELECT line
FROM fortune
WHERE to_tsvector('english', line) @@ to_tsquery('cat & sleep');
                                line
-----------------------------------------------------------------
 People who take cat naps don't usually sleep in a cat's cradle.


SELECT line
FROM fortune
WHERE to_tsvector('english', line) @@ to_tsquery('cat & (sleep | nap)');
                                line
-----------------------------------------------------------------
 People who take cat naps don't usually sleep in a cat's cradle.
 Q:      What is the sound of one cat napping?
```

Word Prefix Search

```
SELECT line
FROM fortune
WHERE to_tsvector('english', line) @@
      to_tsquery('english', 'zip:*')
ORDER BY 1;
                                    line
--------------------------------------------------------------------------
 Bozo is the Brotherhood of Zips and Others.  Bozos are people who band
 … he's the one who's in trouble.  One round from an Uzi can zip
 far I've got two Bics, four Zippos and eighteen books of matches."
 Postmen never die, they just lose their zip.
```

Word Prefix Search

```
EXPLAIN SELECT line
FROM fortune
WHERE to_tsvector('english', line) @@
      to_tsquery('english', 'zip:*')
ORDER BY 1;
                                        QUERY PLAN
----------------------------------------------------------------…
 Sort  (cost=101.21..101.26 rows=21 width=36)
   Sort Key: line
   -> Bitmap Heap Scan on fortune  (cost=24.16..100.75 rows=21 …
        Recheck Cond: (to_tsvector('english'::regconfig, line) …
        -> Bitmap Index Scan on fortune_idx_ts  (cost=0.00..24 …
              Index Cond: (to_tsvector('english'::regconfig, li…
```

```
-- ILIKE is case-insensitive LIKE
SELECT line
FROM fortune
WHERE line ILIKE '%verit%'
ORDER BY 1;
                                    line
------------------------------------------------------------------------
 body.  There hangs from his belt a veritable arsenal of deadly weapons:
 In wine there is truth (In vino veritas).
         Passes wind, water, or out depending upon the severity of the
```

Adjacent Letter Search

```
EXPLAIN SELECT line
FROM fortune
WHERE line ILIKE '%verit%'
ORDER BY 1;
                         QUERY PLAN
-----------------------------------------------------------------
 Sort  (cost=1237.07..1237.08 rows=4 width=36)
   Sort Key: line
   -> Seq Scan on fortune  (cost=0.00..1237.03 rows=4 width=36)
         Filter: (line ~~* '%verit%'::text)
```

```
CREATE EXTENSION pg_trgm;

CREATE INDEX fortune_idx_trgm ON fortune
USING GIN (line gin_trgm_ops);
```

```
SELECT line
FROM fortune
WHERE line ILIKE '%verit%'
ORDER BY 1;
                                    line
----------------------------------------------------------------------
 body.  There hangs from his belt a veritable arsenal of deadly weapons:
 In wine there is truth (In vino veritas).
          Passes wind, water, or out depending upon the severity of the
```

Indexed Adjacent Letters

```
EXPLAIN SELECT line
FROM fortune
WHERE line ILIKE '%verit%'
ORDER BY 1;
                                   QUERY PLAN
---------------------------------------------------------------…
 Sort  (cost=43.05..43.06 rows=4 width=36)
   Sort Key: line
   -> Bitmap Heap Scan on fortune  (cost=28.03..43.01 rows=4 …
         Recheck Cond: (line ˜˜* '%verit%'::text)
         -> Bitmap Index Scan on fortune_idx_trgm  (cost=0.00..…
               Index Cond: (line ˜˜* '%verit%'::text)
```

Word Prefix Search

```
-- ~* is case-insensitive regular expression
SELECT line
FROM fortune
WHERE line ~* '(^|[^a-z])zip'
ORDER BY 1;
                                line
---------------------------------------------------------------------
Bozo is the Brotherhood of Zips and Others.  Bozos are people who band
 … he's the one who's in trouble.  One round from an Uzi can zip
 far I've got two Bics, four Zippos and eighteen books of matches."
 Postmen never die, they just lose their zip.
```

Word Prefix Search

```
EXPLAIN SELECT line
FROM fortune
WHERE line ~* '(^|[^a-z])zip'
ORDER BY 1;
                                        QUERY PLAN
-----------------------------------------------------------------…
 Sort  (cost=27.05..27.06 rows=4 width=36)
   Sort Key: line
   -> Bitmap Heap Scan on fortune  (cost=12.03..27.01 rows=4 …
         Recheck Cond: (line ~* '(^|[^a-z])zip'::text)
         -> Bitmap Index Scan on fortune_idx_trgm  (cost=0.00..…
               Index Cond: (line ~* '(^|[^a-z])zip'::text)
```

Similarity

```
SELECT show_limit();
 show_limit
------------
        0.3

SELECT line, similarity(line, 'So much for the plan')
FROM fortune
WHERE line % 'So much for the plan'
ORDER BY 1;
                          line                     | similarity
---------------------------------------------------+------------
 Oh, it's so much fun,              When the CPU |     0.325
 So much                                           |   0.380952
 There's so much plastic in this culture that      |   0.304348
```

Similarity

```
EXPLAIN SELECT line, similarity(line, 'So much for the plan')
FROM fortune
WHERE line % 'So much for the plan'
ORDER BY 1;
                                QUERY PLAN
----------------------------------------------------------------…
 Sort  (cost=342.80..342.95 rows=59 width=36)
   Sort Key: line
   -> Bitmap Heap Scan on fortune  (cost=172.46..341.06 rows=59 …
         Recheck Cond: (line % 'So much for the plan'::text)
         -> Bitmap Index Scan on fortune_idx_trgm  (cost=0.00..…
              Index Cond: (line % 'So much for the plan'::text)
```

Soundex, metaphone, and levenshtein, and daitch-mokotoff word similarity
comparisons are also available.

Indexes Created in this Section

```
\dt+ fortune
                    List of relations
 Schema | Name    | Type  | Owner    | Size    | Description
--------+---------+-------+----------+---------+-------------
 public | fortune | table | postgres | 4024 kB |

\d fortune and \di+
fortune_idx_text       btree    line                         3480 kB
fortune_idx_lower      btree    lower(line)                  3480 kB
fortune_idx_ops        btree    line text_pattern_ops        3480 kB
fortune_idx_ops_lower  btree    lower(line) text_pattern_ops 3480 kB
fortune_idx_ts         gin      to_tsvector(…)               2056 kB
fortune_idx_trgm       gin      line gin_trgm_ops            4856 kB
```

Use of the Contains Operator @> in this Presentation

```
\do @>
                              List of operators
   Schema   | Name | Left arg type | Right arg type | Result type | Description
------------+------+---------------+----------------+-------------+-------------
 pg_catalog | @>   | aclitem[]     | aclitem        | boolean     | contains
 pg_catalog | @>   | anyarray      | anyarray       | boolean     | contains
 pg_catalog | @>   | anyrange      | anyelement     | boolean     | contains
 pg_catalog | @>   | anyrange      | anyrange       | boolean     | contains
 pg_catalog | @>   | box           | box            | boolean     | contains
 pg_catalog | @>   | box           | point          | boolean     | contains
 pg_catalog | @>   | circle        | circle         | boolean     | contains
 pg_catalog | @>   | circle        | point          | boolean     | contains
 pg_catalog | @>   | jsonb         | jsonb          | boolean     | contains
 pg_catalog | @>   | path          | point          | boolean     | contains
 pg_catalog | @>   | polygon       | point          | boolean     | contains
 pg_catalog | @>   | polygon       | polygon        | boolean     | contains
 pg_catalog | @>   | tsquery       | tsquery        | boolean     | contains
```

https://momjian.us/presentations

https://www.flickr.com/photos/patgaines/